

Homework 3 Report

This is the same as Readme.md

Hardware specs:

Experiments are run on CIMS crunchy1, which according to `lscpu` is x86_64 architecture with 64 cores and each core has 2 threads.

Question 1

(a)

Static scheduling assigns equal size of chunk to the two threads. Therefore, for the first for-loop, the time will be:

thread 1: $1+2+\dots+(n-1)/2,$

thread 2: $(n-1)/2+1 + (n-1)/2+2 + \dots + n-1$

(let's just assume n is an odd number for simplicity, and this does not affect the generality of our discussion.)

Similarly, for the second for-loop, we have:

thread 1: $n-1 + n-2 + \dots + (n-1)/2+1$

thread 2: $(n-1)/2 + (n-1)/2-1 + \dots + 1$

In total, we can see that both of the thread will spend $1 + 2 + \dots + n-1 = n(n-1) / 2$ milliseconds.

After each for-loop, the threads need to synchronize. For the first for-loop, thread 1 spends less time than thread 2, so it will spend:

$$[(n-1)/2+1 + (n-1)/2+2 + \dots + n-1] - [1+2+\dots+(n-1)/2] = (n-1)^2 / 4$$

waiting for thread 2.

Similarly, for the second for-loop, the thread 2 will spend the same amount of time $(n-1)^2 / 4$ waiting for thread 1.

In total, $(n-1)^2 / 2$ will be spend in waiting (synchronization).

(b): For each for-loop, the execution time will be less imbalanced if we use `schedule(static, 1)`.

For the first for-loop:

Thread 1: $1 + 3 + 5 + \dots + (n-2) = (n-1)^2/4$

Thread 2: $2 + 4 + 6 + \dots + (n-1) = (n-1)(n+1) / 4$

For the second for-loop, similarly:

Thread 1: $2 + 4 + 6 + \dots + (n-1) = (n-1)(n+1) / 4$

Thread 2: $1 + 3 + 5 + \dots + (n-2) = (n-1)^2/4$

Combined, the total execution time for both of the thread is the not changed: $1 + 2 + \dots + n-1 = n(n-1) / 2$ milliseconds. But there will be less time spent for waiting to synchronize since it is less imbalanced, so the total running time will be shorter.

(c):

In general, using `schedule(dynamic, 1)` will improve since it will dynamically assign work to threads and results in better load balancing. But in this specific case, `schedule(static, 1)` and `schedule(dynamic, 1)` will be the same. This is because the assignment is not changed given the specific execution time of `f(i)s`.

(d):

Since `f(x)` is an independent function, we can use the directive `nowait` for each of the for-loops. This eliminates the implicit synchronization at the end of each for-loop. So at the end each thread will spend exactly $n(n-1) / 2$ milliseconds and no more waiting time.

Question 2

The max number of threads is set to 64 for this question, and the experiments of thread 1, 2, 4, 8, 16, 32, 64 are shown below:

sequential-scan = 0.638685s			
run with 1 threads	parallel-scan	= 0.534414s	error = 0
run with 2 threads	parallel-scan	= 0.427734s	error = 0
run with 4 threads	parallel-scan	= 0.381150s	error = 0
run with 8 threads	parallel-scan	= 0.297201s	error = 0
run with 16 threads	parallel-scan	= 0.348355s	error = 0
run with 32 threads	parallel-scan	= 0.352518s	error = 0
run with 64 threads	parallel-scan	= 0.352525s	error = 0

Question 3

Since the convergence is slow, I set the number of iterations to 100 for different Ns, which is the same as assignment 1. Note that, to comply with the handing rules, the number of threads, the N and the max iterations have to be modify in the source code. Also, the right hand side `f(x,y)` is set to 1 for simplicity. One should change it to a vector for general case.

For Jacobian 2D:

Threads	seq	1	2	4	8	16	32	64
N=100	0.013248	0.031984	0.017096	0.009588	0.006920	0.007546	0.010835	0.524076
N=1000	1.177591	2.899415	1.478437	0.761617	0.396977	0.204938	0.140522	0.703573
N=10000	130.323083	281.000741	141.376350	70.619156	36.532573	19.125628	12.516912	12.839569

For Gauss-Seidel 2D:

Threads	seq	1	2	4	8	16	32	64
---------	-----	---	---	---	---	----	----	----

Threads	seq	1	2	4	8	16	32	64
N=100	0.015471	0.028812	0.015523	0.008998	0.006812	0.007818	0.013221	0.746888
N=1000	1.402649	2.566101	1.316912	0.686153	0.355611	0.190818	0.144216	1.060069
N=10000	137.948482	240.959659	122.255339	62.039789	32.786162	17.990499	13.334537	13.695364